1.

Checkpoint on slide 11:

```
res0: Array[Int] = Array(1, 2, 3, 4, 5, 6, 7, 8, 9)
```

Checkpoint on slide 55:

```
(SparkCamp,4)
(Spark,3)
(spark,1)
(SparkSQL,1)
(../spark/bin/spark-submit,1)
```

Code for slide 60:

```
val rdd1 = sc.textFile("README.md").filter(_ contains "Spark")
val rdd2 = sc.textFile("spark/docs/contributing-to-spark.md").filter(_ contains "Spar
val wc1 = rdd1.flatMap(l => l.split(" ")).map(w => (w, 1)).reduceByKey(_ + _)
val wc2 = rdd2.flatMap(l => l.split(" ")).map(w => (w, 1)).reduceByKey(_ + _)
val joined = wc1.join(wc2)
```

Checkpoint on slide 60:

```
(Spark,(3,2))
```

2.

The mapper emits one key-value output pair for every input pair of vertices: the output key is the *sorted* vertices and the output value indicates the direction of the edge. The mapper takes directed edge $\langle a, b \rangle$: if $a < b$, it emits $\langle (a, b), 1 \rangle$ and if $a > b$, it emits $\langle (b, a), 2 \rangle$.

The reducer sees all the values $v_1, \ldots, v_n$ for a given key (c, d). If $\{1, 2\} \subseteq \{v_1, \ldots, v_n\}$ then it emits (c, d); otherwise it emits nothing.

No combiner is used; combiners are not likely to help in this problem.

Algorithm 1 Map

---
 function MAP($\langle a, b \rangle$)
  if $a < b$ **then**
   Emit $\langle (a, b), 1 \rangle$
  else
   Emit $\langle (b, a), 2 \rangle$
  end if
 end function

Algorithm 2 Reduce
---
   function REDUCE(Key (c, d), Values $\{v_1, \ldots, v_n\}$)
      if $n < 2$ **then return**
      end if
      for $i = 2, \ldots, n$ **do**
         if $v_i \neq v_{i-1}$ **then**
            Emit $\langle c, d \rangle$
            return
         end if
      end for
   end function

3.

The combiner is the same as the reducer. Let N be the total number of words.

Algorithm 3 Map
---
   function MAP(String s)
      for Word w in s **do**
         Emit $\langle w, 1 \rangle$
      end for
   end function

Algorithm 4 Reduce/Combine
---
   function REDUCE(Key w, Values $\{v_1, \ldots, v_n\}$)
      $s \leftarrow 0$
      for $i = 1, \ldots, n$ **do**
         $s \leftarrow s + v_i$
      end for
      Emit $\langle w, s \rangle$
   end function

Without combiners– The shuffle size is N, and the reduce takes $N \pm O(B)$ operations.

With combiners– After the combine step, there are at most k key-value output pairs, since ther are at most k distinct keys. The shuffle size is kB, and the reduce takes $kB \pm O(B)$ operations.

4.

Let me briefly state the naive solution, which does not parallelize effectively. Run one map-reduce to count the number of elements N. Next, map each input pair $\langle i, a_i \rangle$ to $N-i+1$ output pairs $\langle i, a_i \rangle, \langle i+1, a_{i+1} \rangle, \ldots, \langle N, a_N \rangle$. Reduce by summing all values for a given key. With combiners, the shuffle size is N, and the number of reduce operations is NB where B is the number of mappers. The problem with this solution is that each mapper requires $O(N)$ storage to hold the output keys–but this is on the same order as the size of the entire data.

A better idea is to use divide-and-conquer. The idea is to divide the key-set into B equally-sized partitions: $P_1 = \{1, \ldots, n_1\}, P_2 = \{n_1+1, \ldots, n_2\}, \ldots, P_B = \{n_{B-1}+1, \ldots N\}$. Accordingly define

$$\phi(i) = b \text{ such that } i \in P_b.$$

Then define the partial sums $p_1, \ldots, p_B$ by

$$p_b = \sum_{i \in P_b} a_i,$$

and define quantities

$$u_b = \sum_{c < b} p_b.$$

For $i = 1, \ldots, n$ define the within-partition prefix sums as

$$t_i = \sum_{j \in P_{\phi(i)} : j \leq i} a_j$$

Now observe that

$$s_i = t_i + u_{\phi(i)}$$

Therefore the procedure is as follows

1. (Map/Reduce 1) Count the number of keys N

2. (Map 2) Input: the original key-value pairs. Partition the keys sequentially into B workers

3. (Reduce 2) Each worker $b = 1, \ldots, B$ computes p_b and sends it to the driver

4. The driver computes $u_b = \sum_{c < b} p_c$ and sends u_b to each worker b for $b = 1, \ldots, B$.

5. (Reduce 3) Input: the output of step 2. Each worker $b = 1, \ldots, B$ computes $s_i = u_b + t_i$ and emits $\langle i, s_i \rangle$ for each $i \in P_b$.

To formalize this procedure in the map/reduce framework we have to designate the partition number b as a key throughout steps 2-5. In steps 2 and 5, we have (i, a_i) as values. Note that step 5 uses the same input as step 3.

Algorithm 5 Step 2: Map 2

Parameters $n_1, \ldots, n_{B-1}$ determined in Step 1, and $n_B = N$.
function MAP($\langle i, a \rangle$ from original inputs)
    for $b \in 1, \ldots, B$ **do**
        if $n_b > i$ **then**
            Emit $\langle b, (i, a) \rangle$
            return
        end if
    end for
end function

Algorithm 6 Step 3: Reduce 2
---
function REDUCE(Key b, values (i, a) from step 2)
    $p \leftarrow 0$
    for (i, a) in values **do**
        $p \leftarrow p + a$
    end for
    Emit $\langle b, p \rangle$
end function

Algorithm 7 Step 5: Reduce 3

Parameters $u_1, \ldots u_B$ computed by driver in step 3.
function REDUCE(Key b, values (i, a) from step 2)
    Sort values (i, a) by i
    $s \leftarrow u_b$
    for Value (i, a) in sorted list **do**
        $s \leftarrow s + a$
        Emit $\langle i, s \rangle$
    end for
end function

The cost of the computation is dominated by step 2, when the data is partitioned: this requires a shuffle size of N. This is followed by a reduce in step 3 requiring $O(N/B)$ operations and $O(1)$ space. The driver has to complete $O(B)$ operations in step 4. Finally, each worker has to complete $O(N/B)$ operations in step 5, requiring $O(1)$ memory. The overall number of Map/Reduce iterations is 3, including the initial count.

5.

a. Let A be the adjacency matrix. Then the i, j entry of A^2 is nonzero if and only if there is a path of length 2 from vertex i to j. Therefore the procedure is to evaluate A^2, which costs $O(n^{2.376})$ operations, and check the $n(n-1)/2$ off-diagonals, which costs $O(n^2)$ operations.

b. Let $b = \lfloor n/r \rfloor$, and write $A = [A_1, \ldots, A_b]^T$ where $A_1, \ldots, A_{b-1}$ are $r \times r$ and A_b is $r \times (n - (b-1)r)$. Also write $B = [B_1, \ldots, B_b]$ where $B_1, \ldots, B_{b-1}$ are $r \times r$ and B_b is $r \times (n - (b-1)r)$. Then $C = AB$ is composed of blocks $C_{i,j} = A_i^T B_j$. The cost to compute $C_{i,j}$ is the cost to multiply A_i^T and B_j, which is bounded by $O(r^{2.376})$. Meanwhile, there are b^2 such blocks, where $b^2 = O((n/r)^2)$. Therefore the cost is $O((n/r)^2 r^{2.376})$.

c.

Let $m = n^q$. Then our goal is to provide an algorithm with runtime $O(m^{.55} n^{1.45}) = O(n^{1.45+0.55q})$.

In the proof we will make use of the following Lemma:

Lemma 1 *Let $x = (x_1, \ldots, x_n)$, such that $x \geq 0$, $\sum_i x_i \leq m$, and $\max_i x_i \leq t$. Then*
$$||x||^2 \leq mt$$

Proof. Since $x > 0$, we have $\sum_i x_i = \sum_i |x_i| = ||x||_1$ and also $\max_i x_i = \max_i |x_i| = ||x||_\infty$. Now apply Holder's inequality:
$$||x||^2 \leq ||x||_1 ||x||_\infty \leq mt$$

We first deal with the case $q < 1$. Then there exists at least one node with no edges, hence the graph is not shallow. For any $\epsilon > 0$, it is possible to distinguish the case $q < 1 - \epsilon$ from the case $q \geq 1$ with high probability by computing the proportion of $O(n)$ sampled vertex pairs which have edges. Since $O(n) < O(n^{1.45+0.55q})$ the cost of this statistical test is a non-issue. Henceforth we assume $q \geq 1$.

Having assumed that $q \geq 1$, we can further assume that q is known, since the cost to count the edges is $O(n^2) \leq O(n^{1.45+0.55q})$. We will define

an algorithm with a parameter $p > 0$ depending on q as follows.

We define the algorithm as follows. Order the vertices in descending order by degree. Let $1, \ldots, n_h$ denote all the vertices with degree larger than n^p. Since the number of edges is twice the sum of the degrees, we have $n_h < n^{q-p}$. Let A be the adjacency matrix of the graph, and write $A = [A_1 A_2]$ where A_1 is $n \times n_h$ and A_2 is $n \times (n - n_h)$. By construction, each column of A_2 has at most n^p entries. We will compute

$$A^2 = AA^T = A_1 A_1^T + A_2 A_2^T$$

where A_v is the vth column of A.

To compute $A_1 A_1^T$, apply the result of part b with $r = n_h < n^{q-p}$. Therefore the cost to compute $A_1 A_1^T$ is $O(n^{2+0.376q-0.376p})$.

To compute $A_2 A_2^T$, note that

$$A_2 A_2^T = \sum_{v=n_h+1}^{n} A_v A_v^T$$

Now let S denote a $n \times n$ matrix intialized to 0. For each for each $v = n_h + 1, \ldots, n$, $A_v A_v^T$ is a matrix of zeros and ones, with exactly d_v^2 ones. Therefore, we can compute $A_2 A_2^T$ by iterating over all $v = n_h + 1, \ldots, n$ and incrementing the d_v^2 entries of S which correspond to nonzero entries of $A_v A_v^T$. In this way, the cost of computing $\sum_{v=n_h+1}^{n} A_v A_v^T$ is $\sum_{v=n_h+1}^{n} d_v^2$. Applying the lemma we have $\sum_{v=n_h+1}^{n} d_v^2 \leq 2mn^p$, so the cost of this step is $O(n^{p+q})$.

Now choose $p = 1.45 - 0.45q$. Then the cost of computing $A_1 A_1^T$ is

$$O(n^{2+0.376q-0.376p}) = O(n^{1.45+0.55q})$$

and the cost of computing $A_2 A_2^T$ is

$$O(n^{p+q}) = O(n^{1.45+0.55q})$$

Having computed $A^2 = A_1 A_1^T + A_2 A_2^T$, it remains to check the off-diagonals of A^2. But this takes $O(n^2)$ operations, which is less than $O(n^{1.45+0.55q})$ for any $q \geq 1$.

Hence the cost of the whole procedure is $O(n^{1.45+0.55q}) = O(n^{1.45}m^{0.55})$.